


Classification on hand written numbers

In this activity we will be using MNIST dataset: <https://www.tensorflow.org/datasets/catalog/mnist> (<https://www.tensorflow.org/datasets/catalog/mnist>)

To develop a classifier of hand written numbers as shown in the following figure



This activity is modified from https://keras.io/examples/vision/mnist_convnet/ (https://keras.io/examples/vision/mnist_convnet/)

Title: Simple MNIST convnet

Description: A simple convnet that achieves ~99% test accuracy on MNIST.

Setup

```
In [1]: 1 import numpy as np
        2 from tensorflow import keras
        3 from tensorflow.keras import layers
```

Prepare the data

Model / data parameters

```
In [2]: 1 num_classes = 10
        2 input_shape = (28, 28, 1)
```

the data, split between train and test sets

```
In [3]: 1 (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)
11490434/11490434 [=====] - 1s 0us/step

Scale images to the [0, 1] range

```
In [4]: 1 x_train = x_train.astype("float32") / 255
        2 x_test = x_test.astype("float32") / 255
```

Make sure images have shape (28, 28, 1)

```
In [5]: 1 x_train = np.expand_dims(x_train, -1)
        2 x_test = np.expand_dims(x_test, -1)
        3 print("x_train shape:", x_train.shape)
        4 print(x_train.shape[0], "train samples")
        5 print(x_test.shape[0], "test samples")
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

convert class vectors to binary class matrices

```
In [6]: 1 y_train = keras.utils.to_categorical(y_train, num_classes)
        2 y_test = keras.utils.to_categorical(y_test, num_classes)
```

Build the model

```
In [7]: 1 model = keras.Sequential(
2         [
3             keras.Input(shape=input_shape),
4             layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
5             layers.MaxPooling2D(pool_size=(2, 2)),
6             layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
7             layers.MaxPooling2D(pool_size=(2, 2)),
8             layers.Flatten(),
9             layers.Dropout(0.5),
10            layers.Dense(num_classes, activation="softmax"),
11        ]
12    )
13
14    model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dropout (Dropout)	(None, 1600)	0
dense (Dense)	(None, 10)	16010
=====		
Total params: 34,826		
Trainable params: 34,826		
Non-trainable params: 0		

Train the model

In [8]:

```
1 batch_size = 128
2 epochs = 15
3
4 model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["accuracy"])
5
6 model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, validation_split=0.1)
```

Epoch 1/15
422/422 [=====] - 38s 84ms/step - loss: 0.3719 - accuracy: 0.8880 - val_loss: 0.0853 - val_accuracy: 0.9770

Epoch 2/15
422/422 [=====] - 33s 78ms/step - loss: 0.1141 - accuracy: 0.9654 - val_loss: 0.0591 - val_accuracy: 0.9825

Epoch 3/15
422/422 [=====] - 35s 82ms/step - loss: 0.0885 - accuracy: 0.9720 - val_loss: 0.0513 - val_accuracy: 0.9877

Epoch 4/15
422/422 [=====] - 34s 81ms/step - loss: 0.0730 - accuracy: 0.9778 - val_loss: 0.0507 - val_accuracy: 0.9863

Epoch 5/15
422/422 [=====] - 33s 79ms/step - loss: 0.0643 - accuracy: 0.9801 - val_loss: 0.0383 - val_accuracy: 0.9895

Epoch 6/15
422/422 [=====] - 33s 79ms/step - loss: 0.0577 - accuracy: 0.9816 - val_loss: 0.0367 - val_accuracy: 0.9905

Epoch 7/15
422/422 [=====] - 33s 79ms/step - loss: 0.0524 - accuracy: 0.9836 - val_loss: 0.0350 - val_accuracy: 0.9913

Epoch 8/15
422/422 [=====] - 32s 76ms/step - loss: 0.0480 - accuracy: 0.9843 - val_loss: 0.0350 - val_accuracy: 0.9905

Epoch 9/15
422/422 [=====] - 40s 95ms/step - loss: 0.0459 - accuracy: 0.9847 - val_loss: 0.0349 - val_accuracy: 0.9907

Epoch 10/15
422/422 [=====] - 34s 79ms/step - loss: 0.0438 - accuracy: 0.9855 - val_loss: 0.0313 - val_accuracy: 0.9918

Epoch 11/15
422/422 [=====] - 32s 76ms/step - loss: 0.0413 - accuracy: 0.9864 - val_loss: 0.0330 - val_accuracy: 0.9910

Epoch 12/15
422/422 [=====] - 32s 75ms/step - loss: 0.0368 - accuracy: 0.9882 - val_loss: 0.0290 - val_accuracy: 0.9917

Epoch 13/15
422/422 [=====] - 32s 76ms/step - loss: 0.0374 - accuracy: 0.9879 - val_loss: 0.0295 - val_accuracy: 0.9927

Epoch 14/15
422/422 [=====] - 34s 80ms/step - loss: 0.0363 - accuracy: 0.9877 - val_loss: 0.0311 - val_accuracy: 0.9927

```
uracy: 0.9915  
Epoch 15/15  
422/422 [=====] - 31s 73ms/step - loss: 0.0325 - accuracy: 0.9891 - val_loss: 0.0301 - val_acc  
uracy: 0.9922
```

Out[8]: <keras.callbacks.History at 0x2127a67b890>

Evaluate the trained model

```
In [9]: 1 score = model.evaluate(x_test, y_test, verbose=0)  
        2 print("Test loss:", score[0])  
        3 print("Test accuracy:", score[1])
```

Test loss: 0.02624954842031002

Test accuracy: 0.9909999966621399